Empirical Analysis of Model Structures and Hyper-parameters

We experimented with alternative classification heads and hyperparameter values of ASSORT_S and measured their accuracy. We used the same training data for these experiments, consisting of 2,424 answer posts and 14,165 sentences.

I. CLASSIFICATION HEAD

As Reviewer A suggested, we experimented with different classification heads on top of ASSORT_S's sentence representation, including random forests, decision trees, linear regression, Ada boost, logistic regression, and Naive Bayes Classifier (NBC). Table I shows the model accuracy after replacing the original feedforward neural network (FNN) with each alternative classification head. Column F1 shows the F1 score, and Column Δ shows the model accuracy difference in comparison to the original design (i.e., FNN).

As shown in Table I, using FNN achieved significantly better accuracy, 0.71 in the F1 score. This indicates that FNN has better learnability than other models. In future work, we plan to experiment with more advanced neural network architectures such as LSTM.

TABLE I: Model Accuracy with Different Classification Heads

	F1	Δ
Feedforward NN	**0.71**	-
Random forest	0.56	-0.15
Decision tree	0.54	-0.17
Linear regression	0.65	-0.06
Logistic regression	0.63	-0.08
Ada boost	0.59	-0.12
Naive Bayes Classifier	0.62	-0.09

II. NUMBER OF HIDDEN LAYERS

We experimented with different numbers of hidden layers for the classification head of ASSORT_S. The results are shown in Table II. The results indicate that by adding more hidden layers, the performance of ASSORT_S is not guaranteed to increase. For the simplicity of the model structure and to avoid increasing the risk of overfitting, we choose to include only one hidden layer in the classification head.

Adding one layer to a deep pre-trained language model is a common practice in NLP. For example, in the original BERT paper [1], the authors only used one hidden layer in their classification head. Similarly, the authors in the BERTSum paper [2] also used one hidden layer.

TABLE II: Model Accuracy with Different Numbers of Hidden Layers in the FNN

| Number of hidden layers	Precision	Recall	F1
1 ★	**0.73**	**0.69**	**0.71**
2	0.69	0.71	0.70
3	0.73	0.71	0.72
4	0.73	0.70	0.71

Fig. 1: (a): Learning rate (lr)=1e-4, (b):lr=5e-5, (c):lr=1e-5 (Current design), (d):lr=5e-6

III. LEARNING RATE

We also conducted experiments to decide the learning rate of ASSORT_S. Figure 1 shows the learning curves for different learning rates of ASSORT_S. Our results shows that a learning rate of $1e-5$ gives the best performance.

REFERENCES

[1] J. Devlin, M.-W. Chang *et al.*, "Bert: Pre-training of deep bidirectional transformers for language understanding," *arXiv preprint arXiv:1810.04805*, 2018.

[2] Y. Liu and M. Lapata, "Text summarization with pretrained encoders," *arXiv preprint arXiv:1908.08345*, 2019.